

CRYSTAL CHENG

cjcheng@berkeley.edu | [linkedin.com/in/chengjcrystal](https://www.linkedin.com/in/chengjcrystal) | github.com/chengjcrystal | chengjcrystal.vercel.app

EDUCATION

University of California, Berkeley, Berkeley, CA

Expected Dec 2026

Bachelor of Arts in Computer Science & Applied Mathematics

- Coursework: Artificial Intelligence, Machine Learning, Data Structures, Statistical Modeling, Linear Algebra, Probability Theory
- Involvements: Society of Women Engineers (Treasurer), Web Development at Berkeley (Developer)

WORK EXPERIENCE

Hyve Solutions | Software Engineer Intern

June 2026 - Aug 2026

- Resolved an Azure OpenAI token-limit failure blocking concurrent users by replacing full-catalog retrieval with a Hugging Face embeddings + Neo4j vector-search pipeline, cutting per-request candidates ~95% via a top-150 cap with a keyword fallback.
- Built a full-stack AI test-case search endpoint combining tiered keyword matching, RapidFuzz fuzzy prefiltering, and structured LLM re-ranking constrained to an ID allowlist to eliminate hallucinated results, reused by Analyze's backend to unify retrieval.
- Eliminated duplicate case-detail code across two navigation surfaces by building a single renderer that switches context via a page-state flag, letting the test-plan and test-case managers share one detail view with zero logic duplication.

TrackFly | Software Engineer Intern

Sep 2025 - Dec 2025

- Reduced cross-team integration mismatches on a 15-person team by designing REST API contracts and MySQL schemas that defined a single source of truth for feature data across services.
- Cut per-feature build time by shipping 15+ reusable React/TypeScript components that standardized UI patterns across 10 weeks.

Barobo | Software Engineer Intern

May 2025 - Aug 2025

- Eliminated unauthorized video sharing across all user traffic by engineering single-use AWS CloudFront signed URLs with per-user access control and configurable TTL expiration.
- Scaled to 10,000+ monthly authorized video requests by tracking user access metadata in AWS DynamoDB, enabling reliable per-request authorization without adding infrastructure overhead.
- Cut environment provisioning from 2 hours to 15 minutes by rewriting deployment as reusable Terraform and AWS CLI scripts, making infrastructure reproducible across environments.

Opal Mentorship App | Machine Learning Intern

May 2024 - May 2025

- Raised mentor-mentee match accuracy 12% across cohorts by calibrating 6 scoring weights over a 2-round data-backed process.
- Built a reproducible ML evaluation pipeline in AWS Lambda and DynamoDB that scores compatibility across 4 behavioral dimensions, giving the team versioned, auditable weighting decisions at cohort scale.

PROJECTS

RepoRank Github Search Engine | *Python, FastAPI, SQLAlchemy*

[link](#)

- Served BM25/BM25F search over 157k GitHub repos on a from-scratch field-aware inverted index: p95 23.8ms uncached, sub-millisecond median cached (94.3% hit rate), single-process throughput saturating at ~100 QPS (GIL-bound).
- Built an offline ranking-eval gate (nDCG/MRR/P@k with bootstrap confidence intervals over hand-labeled queries) wired into CI against a frozen 157k index, failing the build when nDCG regresses past a set margin.
- Crawled and indexed 157k repos with a resumable rate-limited GitHub crawler, rebuilding a field-aware inverted index snapshot

FreshCheck Fruit Freshness Classifier | *PyTorch, ONNX, React, Vercel*

[link](#)

- Ran ~14ms fruit-freshness inference fully in-browser with zero backend by training a 24K-parameter CNN in PyTorch and exporting to ONNX (WebAssembly), keeping all images on-device.
- Caught and fixed train/test leakage from duplicate source images by re-splitting on content and perceptual hashes into a leak-free held-out set, reporting honest 97.7% test accuracy (43/44) with a 95% confidence interval.

MBTI Guesser | *Hugging Face Transformers BART-MNLI, DeepFace, OpenCV, Python*

[link](#)

- Built MBTI Guesser (Hugging Face Spaces), a multi-signal personality classifier inferring all four MBTI axes via zero-shot BART-MNLI, and validated the text pipeline on a hand-written unambiguous test set (16/16 correct, >97% mean confidence).
- Late-fused photo (DeepFace + OpenCV) and numeric signals with the text classification using weighted blending that redistributes signal weight to text when photo input is unavailable.
- Diagnosed identity-framing bias in the zero-shot NLI hypothesis template that skewed predictions, rewrote it around evidence-framing language, and added a confidence-gap threshold so ambiguous axes abstain instead of forcing a type.

SKILLS

- Languages: Python, JavaScript, TypeScript, C/C++, SQL, Bash, HTML/CSS (Tailwind)
- ML & AI: PyTorch, TensorFlow, Hugging Face Transformers, ONNX, RAG, Scikit-learn, NumPy, Pandas, OpenCV, DeepFace
- Backend & Data: FastAPI, Node.js, REST APIs, SQLAlchemy, PostgreSQL, DynamoDB, AWS (S3, CloudFront, Lambda)
- Frontend & Systems: React, Next.js, Docker, CI/CD (GitHub Actions, Jenkins), Terraform, Git, Linux, pytest, Vercel